

PLACES 1-6

1



TAPPARA

115

101-135 pts
likely 1st-3rd

2



KOOKOO

107

94-128 pts
likely 2nd-6th

3



LUKKO

105

90-124 pts
likely 3rd-8th

4



ILVES

104

90-124 pts
likely 3rd-8th

5



HIFK

99

83-118 pts
likely 5th-11th

6



JOKERIT

99

84-119 pts
likely 4th-10th

PLACES 7-12

7



SAIPA

98

84-119 pts
likely 4th-10th

8



JYP

97

81-115 pts
likely 5th-11th

9



K-ESPOO

94

82-116 pts
likely 5th-11th

10



KALPA

93

78-113 pts
likely 7th-12th

11



KÄRPÄT

91

76-111 pts
likely 7th-13th

12



ÄSSÄT

89

74-108 pts
likely 9th-14th

PLACES 13-17

13



PELICANS

89

77-111 pts
likely 7th-13th

14



HPK

86

73-107 pts
likely 9th-14th

15



TPS

79

64-98 pts
likely 13th-16th

16



JUKURIT

71

57-90 pts
likely 15th-17th

17



SPORT

69

53-86 pts
likely 15th-17th

HOW IT IS BUILT

1 BUILT FROM PLAYERS, NOT LAST SEASON

Five seasons of goals per game for every player on every roster, with the most recent season weighted heaviest.

2 FOREIGN LEAGUES CONVERTED, NOT GUESSED

An import's scoring is rescaled into Liiga terms — **SHL 1.20, AHL 1.15, Allsvenskan 0.75, Mestis 0.35** — set from players who actually made that move.

3 AGE IS A CLIFF, NOT A SLOPE

Output holds up to about 28, then falls away fast: **0.67** of the previous rate at 35–37, **0.40** at 38 and over. Measured, not assumed.

4 GOALIES CARRY THE DEFENCE

Each team's projected save percentage becomes a goals-against multiplier, pulled toward the league average so one hot 20-game run does not distort it.

5 TWO MODELS, THEN 10,000 SEASONS

A goal-scoring model and a team-strength rating vote **40/60** on every game. The full schedule is then simulated **10,000** times.

THE MATCH MODEL

Goals come from a **Poisson** model, $\lambda = \text{lgAvg} \times \text{off} \times \text{def} \times \text{home_ice}$, with a **Dixon-Coles** correction on the low-score diagonal so the simulated OT rate matches Liiga's observed 23%. Blended **40/60** with a margin-of-victory **Elo** rating ($k=16$).

VALIDATION

Leakage-free backtest over 2023–26, each season rated on prior seasons only: points MAE **12.65**, game log-loss **0.672** vs a **0.686** base rate, standings **Spearman ρ 0.48**. Tuned on log-loss and MAE — ρ is too noisy at $n=4$ to optimise against.

Next: moving the pipeline to scheduled **Snowflake ML jobs** — refitting nightly on real results and re-simulating only unplayed fixtures.